

Siddharth Bhat (20161105)

1 Q1

2 Q2

Consistency: all nodes see the same data at the same time. **Availability:** a guarantee that every request receives a response about whether it was successful or failed. **Partition tolerance:** the system continues to operate despite arbitrary message loss.

The CAP theorem states that you cannot simultaneously have all three; you must make tradeoffs among them. The CAP theorem is sometimes incorrectly described as the choice of picking consistency, availability, and partitioning during **design time**. Really, the theorem allows for databases to choose between these dynamically at **run time**.

When using NoSQL databases, we lose trade-off consistency for a new property called **eventual consistency**. This allows each system to make updates to data and learn of other updates made by other systems within a short period of time, without being totally consistent at all times. For most systems, knowing that it will reach consistency in a short period of time is "good enough", when it allows us to have consistency and availability as well!

3 Q3

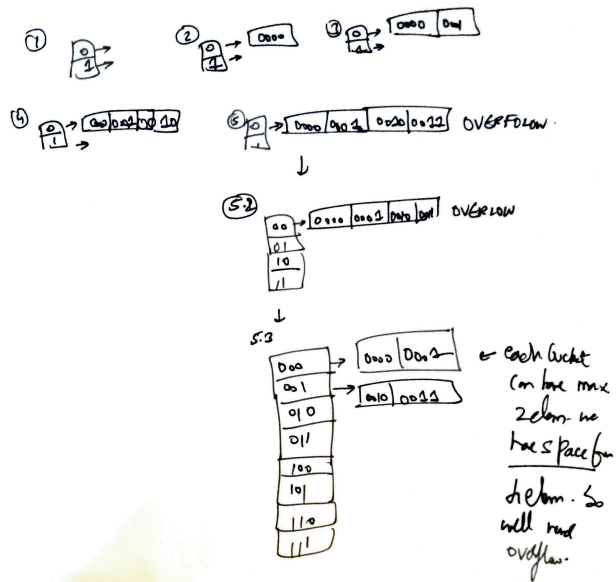
4 Q4

5 Q5

6 Q6

When we insert 0011, there are four records for bucket 0, which overflows. Adding a second bit to the bucket addresses doesn't help, because the first four records all begin with 00. Thus, we go to $i = 3$ and use the first three bits for each bucket address. Now we have space for 4 records in each bucket; however, we can get at maximum 2 records per bucket — we have 4 bit hashes, our buckets are keyed by 3 bits, leaving 1 bit of ambiguity. 1 bit can be 2 records. When 1111 is inserted, we have two records in each bucket.

We assume we have access to MSG.



7 Q7

7.0.1 Q7.1

$SL_x(E)$ for a shared lock by transaction x on elem E . $U_x(E)$ for unlock by transaction x on elem E .

T_1	T_2	T_3
$SL_1(A)$		
$R_1(A)$		
	$SL_2(B)$	
	$R_2(B)$	
		$SL_3(C)$
		$R_3(C)$
$SL_1(B)$		
$R_1(B)$		
	$SL_2(C)$	
	$R_2(C)$	
		$SL_3(A)$
		$R_3(A)$
$XL_1(A)$		
$w_1(A)$		
	$XL_2(B)$	
	$w_2(B)$	
		$XL_3(C)$
		$w_3(C)$

All locks are accepted, since there are no conflicting locks.

7.0.2 Q7.2

Table will be the same as for the above, since we did not have any read action which was followed by a write action of the same element by the same transaction.

7.0.3 Q7.3

8 Q8

$r_1(O_1) \mapsto T_1$ puts $IS(B_1); S(O_1); \text{release}$
 $r_2(O_2) \mapsto T_2$ puts $IS(B_1); S(O_2); \text{release}$
 $r_3(O_1) \mapsto T_3$ puts $IS(B_2); S(O_3); \text{release}$
 $w_1(O_3) \mapsto T_1$ puts $IX(B_2); X(O_3); \text{release}$
 $w_2(O_4) \mapsto T_2$ puts $IX(B_2); X(O_4); \text{release}$
 $w_3(O_5) \mapsto T_3$ puts $IX(B_2); X(O_5); \text{release}$
 $w_1(O_2) \mapsto T_1$ puts $IX(B_2); X(O_3); \text{release}$

9 Q9

9.0.1 9.1

In Undo Logging Logging, we need to write all we need to write all modified data to disk before committing a transaction. This may need a large number of disk I/Os. This is unlike the case of Redo logging, which allows changes to be present in-memory; only need to flush changes before committing.

9.0.2 9.2

Selinger optimization improves upon DP approach by keeping for each not only the plan of least cost, but also plans that have higher cost but produce a result that is sorted in an order that may be useful for parent queries.

9.0.3 9.3

View serializable: If a given schedule is found to be view equivalent to some serial schedule. Alternatively, there are no cycles in the dependency graph.
Conflict serializable: If there are no cycles in the conflict graph.

9.0.4 9.4

We can use strict 2-phase locking for recoverability. This requires that in addition to the lock being 2-Phase, all Exclusive(X) Locks held by the transaction be released until after the Transaction Commits.

9.0.5 9.5

Database operations are in fact relational algebra operations. These operations are pure mathematical expressions, and are generally reads or writes into disjoint pieces of data. This makes them naturally parallelizable.

9.0.6 9.6

File system does not generally have multiple readers and writers to a single file. It also does not need to manage structured data. Hence, many of the ACID like concerns simply do not occur in the case of a file system.

9.0.7 9.7

the commit bit for X is true if and only if the most recent transaction to write X has already committed. The purpose of this bit is to avoid a situation where one transaction T reads data written by another transaction U, and U then aborts. This problem, where T makes a "dirty read" of uncommitted data, certainly can cause the database state to become inconsistent, and any scheduler needs a mechanism to prevent dirty reads.

9.0.8 9.8

- Two-phase locking - 2PL.
- General lock based solutions.
- Timestamp ordering.
- Validation based concurrency control.

Increment based locking is good in this case because it allows to add or subtract a constant from an element, which is what most kinds of bank transactions are. Increment locks on the same element do not conflict with each other.

9.0.9 9.9

recovery manager will have to DODO

9.0.10 9.10

all trees of n vertices is n^{n-2} . Number of left-deep trees is $n!$. n^{n-2} is much larger than $n!$.