

# Siddharth Bhat

---

## Skills

### Formally Verified Decision Procedures for Compiler IRs:

- Top 20 contributors to the Lean theorem prover overall.
- #2 contributor to Lean's bitvector theory.
- Coauthor of the Lean4 metaprogramming book.
- Published new, formally verified decision procedures for parametric bitvector theory.
- Led the lean-mlir project, for verifying optimizations in MLIR.
- Given >4 talks @ LLVM dev meeting (formal semantics for LLVM & MLIR).

### Loop Optimization & HPC:

- Worked on Polly, LLVM's polyhedral loop optimizer.
- 121 commits in Polly, and the #3 contributor to the Polly. Added support for GPU code generation.
- Mentored students for LLVM in Google Summer of Code (GSoC 2016).

### Compilers for Functional Languages:

- #2 contributor to Asterius, a Haskell → WebAssembly compiler.
- I implemented a Haskell-style runtime on top of the WASM runtime. Contributions merged into Asterius, and eventually into GHC proper.
- Primary author of Lean's LLVM backend.
- Papers on designing MLIR based IRs for quantum & functional compilation.

### AI for Maths:

- RL for proofs in the F\* proof assistant @ Microsoft Research ("Towards neural synthesis for SMT-assisted proof-oriented programming" @ICSE 2025, best paper).
- Established stronger baselines for AI based geometric theorem proving.

## Education

PhD **University of Cambridge**.

(2024 - Ongoing)

PhD **University of Edinburgh (moved to Cambridge)**.

(2022 - 2024)

Master by **International Institute of Information Technology Hyderabad India**.

Research (2020 - 2021)

B.Tech **International Institute of Information Technology Hyderabad India**.

2015 - 2020

## Publications

Certified Decision Procedures for Width-Independent Bitvector Predicates: **Siddharth Bhat (1st)**, Léo Stefanesco, Chris Hughes, Tobias Grosser. OOPSLA 2025

Interactive Bit Vector Reasoning using Verified Bitblasting: Henrik Böving, *Siddharth Bhat*, Alex Keizer, Luisa Cicolini, Leon Frenot, Abdalrhman Mohamed, Léo Stefanesco, Harun Khan, Josh Clune, Clark Barrett, Tobias Grosser. OOPSLA 2025

Verifying Peephole Rewriting in SSA Compiler IRs: **Siddharth Bhat (1st)**, Alex Keizer, Chris Hughes, Andres Goens, Tobias Grosser. ITP 2024

Verifying Wu's Method can Boost Symbolic AI to Rival Silver Medalists and AlphaGeometry to Outperform Gold Medalists at IMO Geometry: Shiven Sinha, Ameya Prabhu, Ponnurangam Kumaraguru, *Siddharth Bhat*, Matthias Bethge. NeurIPS 2024 Workshop MATH-AI

Towards Neural Synthesis for SMT-Assisted Proof-Oriented Programming: Saikat Chakraborty, Gabriel Ebner, *Siddharth Bhat*, Sarah Fakhouri, Sakina Fatima, Shuvendu Lahiri, Nikhil Swamy. ICSE 2024

Rewriting Optimization Problems into Disciplined Convex Programming Form: Ramon Fernandez Mir, *Siddharth Bhat*, Andres Goens, Tobias Grosser. CICM 2024

Guided Equality Saturation: Thomas Koehler, Andres Goens, *Siddharth Bhat*, Tobias Grosser, Phil Trinder, Michel Steuwer. POPL 2024

Lambda the Ultimate SSA: **Siddharth Bhat (1st)**, Tobias Grosser. CGO 2022

QSSA: An SSA based IR for Quantum Computing: Anurudh Peduri, *Siddharth Bhat*, Tobias Grosser. CC 2021

Optimizing Geometric Multigrid Computation using a DSL Approach: Vinay Vasista, Kumudha KN, **Siddharth Bhat**, Uday Bondhugula. Supercomputing (SC), Nov 2017

Word Embeddings as Tuples of Feature Probabilities: **Siddharth Bhat (1st)**, Alok Debnath, Souvik Banerjee, Manish Shrivastava Representation Learning for NLP, 2020

## Internship Experience

- Sep-Nov '24 **Amazon Web Services, Automated Reasoning Group, Austin.**  
Deciding memory (non)interference in `linsym`, a Lean-based ARM symbolic simulator
- Jul-Sep '23 **Microsoft Research, Redmond.**  
Retrieval Augmented theorem proving for the Fstar proof assistant.
- July 1-10 '23 **Adjoint School, Glasgow.**  
Researched Markov categories and their relationship to probabilistic programming.
- May-Jul '19 **Intern at Tweag.io, Paris, France.**  
Re-implemented portions of GHC(Glasgow Haskell Compiler) runtime for [Asterius](#) ([link](#)), a Haskell to WebAssembly compiler. Involved Haskell, C, and WebAssembly.
- Summer 2018 **Visiting research intern at ETH Zurich, Zurich, Switzerland.**  
Investigating formal verification of polyhedral compilation. [PolyIR](#) ([Link](#)) is a formal specification of polyhedral programs.
- Summer 2018 **GSoC mentor, Polly Labs.**  
Mentoring a project to enable Polly's loop optimisations into Chapel.
- Mar-Dec '17 **ETH Zurich, Research Intern at SPCL, Zurich, Switzerland.**  
Worked on Polly, a polyhedral loop optimizer for LLVM.
- May-Jul '16 **Research Intern, IISc Bangalore, Bangalore.**  
Worked on PolyMage, DSL compiler for optimising loop transforms. Contributed to ISL and PLUTO. Implemented tiling patterns, optimised PolyMage for stencils.
- Summer 2016 **Selected for GSoC 2016, Google.**  
Binding SymEngine, a symbolic math library to Haskell. Had to drop this to intern at IISc, Bangalore. Still maintain the library (`symengine.hs`)

Summer 2015 **GSoC 2015, Google.**

Worked on VisPy, a pure Python graphics library which uses OpenGL internally for performance. Successfully completed.

## Open Source Contributions

**Lean4** Co-developed the bitblasting theory for Lean's bitvector automation, wrote the LLVM backend for the compiler.

**Rocq** Submitted issues, bug-fixes, helped improve developer documentation.

**VE-LLVM** Collaboration with VE-LLVM, a formal semantics of the LLVM compiler toolchain in Coq

**Polly** Implementing support for Fortran, added unified memory abilities to the CUDA backend within Polly, a polyhedral loop optimiser for LLVM. ([Link to commits](#))

**Symengine.hs** GSoC 2016. Haskell bindings to SymEngine, a C++ symbolic manipulation library.

**VisPy** GSoC 2015. Rewrote scene graph for performance. Added visuals, high level API for easy use of plotting. Implemented auto-resizing with **Cassowary**, a linear optimisation library.

**Rust** Contributed to the Rust compiler and ecosystem. Found compiler errors, fixed libraries. Was part of **Piston**, group of Rust programmers who experimented with writing game engines.

**Haskell** Contributed to the Haskell ecosystem. Reported and fixed bugs in *stack*, *stackage*, *diagrams*, *GHC*, etc. ([Link to GHC commits](#)).

**PLUTO** Source to Source C optimiser for loop nests. Improved the PLUTO API that had gone out of sync with master. Discovered bugs in PLUTO for diamond tiling transforms

**PolyMage** DSL compiler that generates C code. Uses **Polyhedral Compilation** Extended the compiler to add stencils, time iterated-stencils.

**PPSSPP** C++ open source PSP emulator. Wrote most of the touch handling code. Implemented atomic locks for audio performance.

## My Projects

**Lean-MLIR** Formal semantics for the MLIR compiler framework, defined within the Lean4 proof assistant.

**Iz** An MLIR based compiler backend for the Lean4 proof assistant.

**Lean4 Metaprogramming Book** A textbook on metaprogramming in Lean4. I wrote the chapters on tactics and metaprogramming for embedded DSLs.

**Lean-to** Jupyter kernel for the Lean4 proof assistant.

**Simplexhc** A custom compiler for a subset of Haskell. The goal is to try and apply *Polyhedral compilation* ideas to compile a lazy, pure, functional programming language with LLVM as a backend. Has **64 stars** on github.

**Sublime Bookmarks** A plugin for sublime text to quickly jump between pieces of your codebase. **26k downloads** and counting.

**Cellular Automata** A collection of Cellular Automata written in Haskell. Uses **Comonads** for abstraction. **130 stars** on Github.

**Teleport** A simple tool to switch between projects written in Haskell. Shows how to write "real world Haskell". Published as a **Literal Haskell tutorial**. **90 stars** on github

**TIMi** A visual interpreter of the **template instantiation machine** to understand evaluation of lazy functional languages. **51 stars** on github.

## Miscellaneous

**Barvinok** Talk at ETH Zurich: Slides describing the Barvinok algorithm to count lattice points in polyhedra

**FunctionalConf '19** Talk on implementing embedded probabilistic programming languages in Haskell ([Slides](#))

**Haskell Exchange 2020** Talk on optimizing `smallpt-hs` (a port of a raytracer to haskell) to beat C++ performance ([Slides](#))

**FPIIndia** Talk on egg: fast and extensible equality saturation. ([Slides](#))

**Theory seminar, winter '19** Talk on impossibility of compass-straightedge constructions using field theory.

**math.se** Answer on **math.stackexchange**. **8312 reputation, top 4% overall**. Abstract algebra and differential/algebraic geometry.

## Talks & Presentations

**Euro LLVM Dev 2025:** How to trust your peephole rewrites: automatically verifying them for arbitrary width!. **US LLVM Dev 2024:** lean-mlir: A workbench for formally verifying peephole optimizations in MLIR. **US LLVM Dev 2023:** (Correctly) Extending dominance to MLIR Regions. **US LLVM Dev 2023:** MLIR Side Effect Modelling. **Euro LLVM Dev 2022:** MLIR for Functional Programming. **FPIIndia 2021:** Equality Saturation. **Functiona Conf 2019:** Monad-bayes: Probabilistic programming in Haskell.

## Awards

<https://www.renaissancephilanthropy.org/mathbench-towards-evaluating-natural-language-proofs>

One of 30 research groups awarded out of over 280 applicants.